

F6. Evidence N-Ladder

Evaluation units are separated to avoid treating clustered model assessments as independent calls.

Test split

Unit: audio rows

N = 258

ASR model comparison

Selected provenance

Unit: candidate rows / outputs

N = 300

selection and provenance

Human-reviewed audit

Unit: audio rows

N = 30

decision-critical review unit

Reviewed assessments

Unit: model-row assessments

N = 90

predictor and recovery replay

Cluster rule

The 90 reviewed assessments are clustered within 30 audio rows; uncertainty should use row-clustered bootstrap or leave-one-row-out sensitivity.