

When Low WER Becomes Dangerous: Counterfactual Decision-Stability ASR for High-Stakes Speech-Driven Decision Systems

2026-05-28

Abstract

Speech-driven decision systems increasingly transform conversations into operational signals for routing, escalation, compliance monitoring, and case-handling decisions. In high-stakes anti-fraud calls, small ASR differences can land on decision atoms such as negation, amount, actor, action, time, and intent. A transcript can therefore remain close under WER or CER while a plausible alternative transcript changes the downstream decision.

We propose Counterfactual Decision-Stability ASR (CDS-ASR), a framework that evaluates whether downstream decisions remain stable under acoustically and semantically plausible transcript alternatives. CDS-ASR extracts risk atoms, constructs plausible ASR alternatives, scores decision instability with Counterfactual Escalation Instability Score (CEIS), and applies constrained recovery or conservative machine action for high-risk spans while preserving aggregate-only reproducibility for sensitive call evidence.

Across a six-model 258-row split comparison, selected-300 high-stakes provenance outputs, and a 30-row/90-model-assessment human-reviewed audit, CDS-ASR provides a consequence-centered evidence layer beyond transcript similarity. In the scoped selected-300 human-reviewed audit, CEIS achieves the highest decision-change AUC and a zero-false-negative operating point. In aggregate policy replay, risk-triggered policies, including CEIS-triggered conservative action, eliminate high-risk missed and critical miss counts under the aggregate-only claim boundary. These results support decision-stability evaluation as a scoped companion to transcript accuracy, semantic metrics, and confidence-aware correction.

Introduction

Speech is no longer only a record of what happened in a call. In modern contact centers, speech is increasingly converted into operational signals: categories, summaries, compliance cues, real-time alerts, routing decisions, and escalation

support. Commercial contact-center analytics already describe voice conversations and transcripts as inputs for automated categorization, summarization, compliance monitoring, and real-time assistance (Amazon Web Services 2026a, 2026b). This turns ASR output into a decision substrate. A transcript can influence how a case is categorized, which queue receives it, whether an alert is raised, and whether a human reviewer ever sees the full call context.

Anti-fraud calls make this speech-to-decision pipeline concrete. Taiwan’s National Police Agency describes the 165 hotline as a public anti-fraud channel where staff record incident details and provide information to victims (National Police Agency, Ministry of the Interior 2024). The operational scale is large: the FBI’s 2025 Internet Crime Report context reported more than one million IC3 complaints and large cyber-enabled fraud losses (Federal Bureau of Investigation 2026b, 2026a). In this setting, the important spoken facts are often small: whether money was already transferred, whether the amount was 30,000 or 300,000, who initiated contact, when the event happened, whether the caller is certain, and whether the scenario matches a known scam pattern. These are not merely words. They are decision atoms.

This paper starts from the following practical risk: a transcript can be linguistically close and still operationally unsafe. A single ASR ambiguity around negation, amount, actor, action, time, intent, or scam pattern can move a case from routine review to priority review, or from escalation to unsafe downrouting. In this paper, “dangerous” is operationalized as ASR-induced decision change that can produce unsafe downrouting, high-risk missed escalation, or critical miss under the declared anti-fraud triage policy. The title is therefore a measurable claim about decision consequences, not a general safety slogan.

Existing ASR evaluation and repair methods provide the foundation for this problem. WER and CER remain necessary because they make transcript accuracy auditable and comparable. Semantic ASR metrics make the evidence more task-aware: Kim et al. motivate Semantic Distance by showing that WER can fail to reflect downstream spoken-language-understanding behavior, and Rugayan et al. evaluate Aligned Semantic Distance as a semantic metric that better aligns with human judgments and downstream tasks (Kim et al. 2021; Rugayan, Salvi, and Svendsen 2023). Transcript repair also advances the pipeline. Naderi et al. study post-hoc ASR correction with LLMs and use confidence-based filtering to reduce harmful changes to likely accurate transcripts (Naderi et al. 2024). Selective prediction, reject-option, and conformal prediction work further show how systems can trade coverage, abstention, and uncertainty when acting on uncertain predictions is costly (Chow 1970; Geifman and El-Yaniv 2017, 2019; Angelopoulos and Bates 2021).

These methods make speech systems more measurable, semantic, correctable, and uncertainty-aware. The remaining gap is the decision-stability question. A semantic metric can report that two transcripts are close or far. A correction model can make a transcript more fluent or more likely. A confidence filter can avoid some risky changes. Yet a high-stakes workflow also needs to know

whether plausible ASR alternatives would change the downstream action. In anti-fraud triage, the actionable question is not only “How accurate is the transcript?” but also: “Would another plausible transcript lead to a different safe action?”

We propose Counterfactual Decision-Stability ASR (CDS-ASR) to answer that question directly. CDS-ASR extracts decision-critical risk atoms, generates acoustically and semantically plausible transcript alternatives, measures decision instability with Counterfactual Escalation Instability Score (CEIS), and evaluates constrained recovery or conservative machine action for high-risk spans. WER, CER, semantic metrics, SRES, confidence, and correction baselines remain comparison layers. The central contribution is the new evaluation view: ASR for high-stakes speech-driven systems should be judged by whether the downstream decision remains stable under plausible transcript alternatives.

Verified source anchors for this introduction are maintained in `citation_seed.md` and `references.bib`, with access dates refreshed on 2026-05-28. Empirical manuscript claims point to the aggregate artifacts listed in the Appendix.

Related Work

Transcript-Centered ASR Metrics

Transcript accuracy metrics provide the reproducible baseline for ASR comparison. WER and CER remain necessary because they allow stable comparison across models and splits when tokenization, normalization, micro/macro scope, and zero-reference handling are declared. This manuscript keeps WER/CER as auditable surface metrics and uses the repo’s metric audit to define the Chinese ASR reporting policy. The paper does not discard transcript accuracy; it treats WER/CER as the auditable surface layer that must be reported before stronger downstream claims are made.

Semantic And Downstream-Aware ASR Evaluation

Semantic ASR metrics extend transcript evaluation toward meaning preservation and downstream task behavior. Kim et al. introduce Semantic Distance for ASR performance analysis toward spoken-language understanding, motivating the move from literal correctness to semantic correctness for intent recognition, slot filling, semantic parsing, and named entity recognition (Kim et al. 2021). Rugayan et al. evaluate Aligned Semantic Distance against WER and show its value for human perception and downstream NLP tasks (Rugayan, Salvi, and Svendsen 2023). CDS-ASR uses this line of work as a foundation and adds a direct decision-stability target: whether plausible ASR alternatives change a high-stakes downstream decision.

Transcript Repair And Conservative Decision Action

LLM and confidence-aware correction methods provide transcript-repair baselines. Naderi et al. show that confidence-based filtering can guide post-hoc LLM correction of ASR transcripts and reduce the chance of changing likely accurate transcripts (Naderi et al. 2024). CDS-ASR complements this line by evaluating the decision interval that remains when plausible ASR alternatives affect risk atoms. Recovery in this manuscript is automatic and machine-bounded: human review supplies evaluation labels and governance evidence, while risk-triggered policies supply the scoped intervention test.

The conservative-action layer also connects to selective prediction and rejection-option research. Chow formalizes the error-reject tradeoff for recognition systems (Chow 1970), while modern selective classification and SelectiveNet show how neural systems can abstain or reject to manage risk-coverage tradeoffs (Geifman and El-Yaniv 2017, 2019). Conformal prediction further motivates distribution-free uncertainty sets for high-risk systems (Angelopoulos and Bates 2021). CDS-ASR applies this governance intuition to speech-driven decision evidence: when transcript alternatives affect decision atoms, the system should expose the instability or take conservative action rather than treat one transcript as fully stable.

High-stakes speech domains reinforce the need for consequence-aware ASR evaluation. In psychotherapy, Miner et al. show that ASR feasibility depends on the clinical use case and that individual-level safety monitoring requires a more cautious threshold than population-level language analysis (Miner et al. 2020). CDS-ASR brings the same claim-evidence discipline to anti-fraud calls by separating transcript accuracy, semantic risk, decision instability, and aggregate recovery evidence.

This paper positions ASR models as evidence-producing components. The protagonist is the downstream decision that changes when an acoustically plausible transcript difference lands on a decision-critical atom.

Method

Problem Formulation

We study speech-driven decision systems where an audio segment is transcribed by ASR and then used by a downstream escalation workflow. The downstream label space includes routine review and higher-priority escalation states. The central unit is a model-sample: one ASR hypothesis for one audio row, evaluated against aggregate risk and decision-stability evidence.

The paper asks whether the decision remains stable under plausible transcript alternatives. This makes transcript accuracy a supporting layer and decision-stability the primary high-stakes safety target.

CDS-ASR Pipeline

CDS-ASR evaluates high-stakes ASR through a decision-stability pipeline:

```
audio
-> ASR transcript and runtime/quality signals
-> risk atom extraction
-> plausible counterfactual transcript variants
-> downstream decision model
-> SRES and CEIS scoring
-> constrained recovery or conservative machine action
```

Risk atoms are transcript spans whose alternatives can affect downstream decisions. This manuscript focuses on negation, amount, action, actor, time, intent, uncertainty, and scam-pattern atoms.

CEIS scores the maximum decision-flip risk among plausible variants. In this paper-facing implementation, the plausibility term is a bounded proxy rather than a calibrated acoustic posterior:

```
CEIS(x) = max over v in V(x) [
    Plausibility(v | x) * RiskAtomWeight(v) * DecisionDistance(f(x), f(v))
]
```

`Plausibility(v | x)` denotes a bounded proxy plausibility score derived from model disagreement, Mandarin phonetic ambiguity, domain-slot alternatives, and available ASR/runtime signals. `RiskAtomWeight(v)` is a risk-atom class weight defined by the decision-critical atom schema. `DecisionDistance` maps the downstream label space into an ordinal safety distance over `no_escalation`, `review`, `priority_review`, `critical_escalation`, conservative machine action, and abstention, with critical misses treated as the highest-risk direction for conservative action.

Downstream Decision Function

The downstream decision function `f` is fixed before CEIS scoring. It maps each transcript or plausible variant into a declared triage action:

```
f(transcript) in {
    no_escalation,
    manual_review,
    priority_review,
    critical_escalation,
    conservative_machine_action,
    abstain
}
```

The distance function is policy-aligned rather than language-model-generated: same-action pairs have distance 0, neighboring review/escalation states have distance 1, and movement from no escalation or routine handling to

critical escalation receives the largest distance. Unsafe downrouting of a critical event receives the maximum penalty used by the policy replay. The method contract is versioned in `docs/downstream_decision_contract.md` and `80_semantic_risk_asr/scoring/ceis_config.json` so CEIS rankings can be reconstructed from aggregate artifacts.

CEIS Scale Discipline

All CEIS components are treated as bounded comparison terms, not calibrated probabilities. `Plausibility(v | x)` is a proxy plausibility score; the paper does not claim an acoustic posterior unless a later implementation supplies one. `RiskAtomWeight(v)` and `DecisionDistance(f(x), f(v))` must be normalized or tabulated before final submission. The next method-hardening analyses are: uniform-weight ablation, no-plausibility ablation, binary-decision-flip ablation, max versus top-k mean aggregation, and CEIS behavior by atom class. Those analyses can be added only from aggregate-safe recomputation or reported as planned validation if the current artifact package cannot reconstruct them.

Thresholds reported in the current predictor table are diagnostic operating points selected on the scoped reviewed audit. They are useful for aggregate comparison and reviewer inspection, but they are not frozen deployment thresholds unless later selected on a separate development set.

Recovery remains automatic and machine-bounded. Human review is used as an evaluation and governance layer, not as the proposed recovery method. The recovery policies compare no recovery, confidence-only trigger, SRES-triggered recovery, CEIS-triggered conservative action, and CEIS ensemble arbitration.

Risk Atom Schema

The risk atom schema focuses on decision-critical spans:

Risk atom	Decision role	Example effect
Negation	reverses event state	transfer vs no transfer
Amount	changes severity	30,000 vs 300,000
Action	changes case stage	asking vs reporting vs transferring
Actor	changes scam interpretation	bank, police, family, customer service
Time	changes urgency	today, yesterday, next week
Intent	changes routing	inquiry, report, complaint
Uncertainty	changes confidence and safe action	sure, maybe, not sure
Scam pattern	changes case type	investment, fake police, installment cancellation

Counterfactual Variant Contract

Counterfactual variants are plausible ASR alternatives concentrated around risk atoms. They are not generic paraphrases. A variant can be supported by

acoustic ambiguity, Mandarin phonetic confusion, domain-slot alternatives, or model disagreement. The paper-facing claim is not that every possible variant is enumerated; it is that the generated alternatives make downstream decision-instability measurable.

Recovery Policy Contract

The five-policy recovery experiment is:

1. no recovery;
2. confidence-only trigger;
3. SRES-triggered recovery;
4. CEIS-triggered conservative action;
5. CEIS ensemble arbitration.

The scoped recovery claim is evaluated only on aggregate human-reviewed selected-300 evidence. Per-row details stay local or ignored. At the policy layer, both SRES-triggered recovery and CEIS-triggered conservative action are reported as risk-triggered conservative policies; CEIS’s distinct contribution is evaluated primarily in the predictor layer and in its decision-stability framing.

Reproducibility Layer

The paper-facing ASR surface metric is `cer_zh_micro`. The supplemental word metric is `wer_zh_jieba_micro`. Raw whitespace WER and legacy stored WER fields are audit-only because they are unstable for unsegmented Chinese transcripts.

The metric policy is validated by `70_experiments/runs/wer_metric_audit_2026_05_25/`, which records manifest checks, zero-reference-unit checks, tokenizer and normalization policy, and `jiwer` cross-checks.

The target transcription locale is Taiwan Traditional Chinese. New candidate models must satisfy the strict Taiwan Traditional Chinese locale gate before promotion. Previously completed comparable baselines are retained only as disclosed baselines, with locale-violation counts reported and no promotion claim attached. Post-decode OpenCC or other conversion can be evaluated only as a deployment repair lane; it cannot be used to claim that the raw ASR model passed the locale gate.

Selected-300 human audit evidence is tracked only through aggregate outputs. The local transcript-bearing audit sheet remains ignored. The reviewed scope is risk atoms, decision-change labels, expected safe action, confidence, per-model assessment fields, and per-row timing. The current aggregate status is `review_complete` with 30/30 reviewed rows and 90/90 reviewed model assessments.

All paper-facing claims point to aggregate artifacts. The manuscript does not require raw audio, raw transcripts, selected row IDs, hypotheses, reviewer notes,

or model weights.

Evaluation Units And N-Ladder

The manuscript separates evidence units explicitly so that model assessments are not mistaken for independent audio rows.

Layer	Unit	N	Role
Test split	audio rows	258	ASR model comparison
Selected high-stakes provenance	selected candidate rows / outputs	300	row-selection and provenance
Human-reviewed audit rows	audio rows	30	decision-critical review unit
Human-reviewed model assessments	model-row assessments	90	predictor and recovery evaluation

Because the 90 model assessments are clustered within 30 audio rows, inferential uncertainty should be estimated with row-clustered bootstrap resampling or leave-one-row-out sensitivity analysis. Point estimates are reported as scoped descriptive evidence; deployment claims do not treat the 90 assessments as independent calls.

Selection Provenance And Enrichment

The selected-300 audit surface is an enriched high-stakes sample, not a prevalence-preserving sample of all anti-fraud calls. The selection summary uses aggregate-safe provenance signals from SRES scoring, CEIS scoring, and downstream escalation decisions, including high-risk missed, critical miss, unsafe downrouting, model disagreement, high proxy risk, risk-score fill, and clean-control strata. The aggregate selection thresholds recorded for the current package are `low_wer_threshold=10.0`, `sres_threshold=20.0`, and `ceis_threshold=5.0`.

Predictor results therefore support metric behavior on an enriched high-stakes audit surface rather than population-level risk prevalence. If a submission emphasizes CEIS/SRES separation, a sensitivity check should exclude rows directly selected by CEIS or SRES family signals and confirm whether the same directional pattern remains.

Variant Coverage And Human Review Reliability

The current aggregate-only coverage audit reports CEIS top-atom proxy coverage over the 90 reviewed model assessments. It records 90 aggregate proxy observations across 30 reviewed rows, with top-atom counts of negation 47, amount 37, action 3, and actor 3. The reviewed selection surface also covers risk-signal atoms in aggregate: action 25, actor 15, amount 23, negation 14, and scam pattern 23 selected rows. These counts support submission-safe

coverage auditing, not release of variant text. Source-specific coverage is available for model-disagreement provenance in aggregate; phonetic, domain-slot, runtime-signal, rejected-variant, and full variant-generation logs are not currently reconstructed as release artifacts. Stronger generator claims therefore require a future aggregate variant-generation log.

Human labels come from a completed expert audit rather than multi-annotator adjudication. A second-reviewer blinded spot-check is not included in the current submission-prep package, so inter-annotator agreement is not claimed. The claim boundary is therefore a completed single-expert audit over 30 reviewed rows and 90 model assessments. A future validation extension can add a blinded 5-10 row spot-check while preserving the frozen selected-300 review boundary.

Experiments

Experiment 1 evaluates comparable ASR model evidence on the canonical 258-row test split. This table supports model-comparison and split evidence, not the paper-grade selected-300 risk/recovery claims.

Experiment 2 uses selected-300 high-stakes proxy outputs as input provenance. The selected-300 proxy outputs identify an enriched high-stakes evidence surface and the review sample, but their raw rows and transcript-bearing metric inputs remain local or ignored. This experiment supports selection provenance, not population prevalence.

Experiment 3 evaluates WER, CER, SRES, and CEIS against human-reviewed decision-change labels over 90 model assessments from the selected-300 audit. These assessments are clustered within 30 reviewed audio rows. This is the paper-grade predictor evidence for metric-insufficiency and decision-stability claims, reported as scoped descriptive evidence until row-clustered uncertainty is added.

Experiment 4 evaluates five recovery policies against the same human-reviewed evidence layer. This is aggregate policy replay evidence, not a live causal deployment trial.

Candidate ASR and multimodal models are kept in a separate exploratory lane. They must not enter the main ASR table until they pass field-contract, runtime-validity, and Taiwan Traditional Chinese locale gates in order.

Results

The Results section separates comparable ASR evidence, candidate-lane boundaries, human-reviewed predictor evidence, and human-reviewed recovery evidence. This separation preserves claim-evidence alignment: the main ASR benchmark supports model-comparison context, while selected-300 human-reviewed outputs support paper-grade risk and recovery claims.

Human review supplies evaluation labels only. The proposed recovery policies are automatic, aggregate-evaluated policies; no transcript-bearing row content is required for paper-facing claims.

Table 1. Main ASR Benchmark Table

Table 1 reports six comparable ASR runs on the canonical 258-row split. The paper-facing primary ASR metric is `cer_zh_micro`; `wer_zh_jieba_micro` is reported as the supplemental segmented word metric. Unsafe downrouting and high-risk missed counts are scope-controlled split/model-comparison evidence, not the final selected-300 human-reviewed risk claim.

Only models with completed comparable 258-row split evidence are included.

Run	Model family	Rows	<code>cer_zh_micro</code>	<code>wer_zh_jieba_micro</code>	Unsafe down-routing	High-risk missed	Locale violation rows	Paper use
<code>breeze_asr25_partial_encoder_legacy_best_test_split</code>	Breeze-ASR-25 partial encoder	258	15.04	21.53	7	4	0	Scope-controlled split/model comparison
<code>breeze_asr25_lora_legacy_best_test_split</code>	Breeze-ASR-25 LoRA	258	18.23	25.59	10	7	0	Scope-controlled split/model comparison
<code>breeze_asr25_base_test_split</code>	Breeze-ASR-25 base	258	22.72	30.39	34	30	0	Scope-controlled split/model comparison
<code>breeze_asr26_test_split</code>	Breeze-ASR-26	258	24.27	32.29	27	22	0	Optional dialect-aware comparator
<code>whisper_large_v2_test_split</code>	Whisper large-v2	258	24.72	32.23	33	28	1	Comparable ASR baseline
<code>whisper_small_test_split</code>	Whisper small	258	34.86	43.44	76	70	4	Comparable ASR baseline

Evidence source: `70_experiments/runs/janus_258_test_split_asr_cds_proxy/asr_cds_proxy_comparison`

The legacy Breeze-ASR-25 partial encoder is the strongest current hypothesis generator on the 258-row split, with the lowest `cer_zh_micro` and the lowest unsafe downrouting and high-risk missed counts among the six comparable runs. This supports using it as a strong ASR evidence layer while keeping final risk and recovery claims tied to the human-reviewed selected-300 outputs.

Table 2. Candidate And Exploratory Lane

Table 2 reports bounded candidate-lane evidence for models that are not promoted to the main 258-row or selected-300 tables. These gates document feasibility, locale behavior, and runtime constraints without mixing exploratory candidates with comparable ASR baselines.

These models are bounded negative, feasibility, or runtime evidence. They do not enter the main ASR benchmark table.

Candidate	Current gate	Rows	cer_zh	wer_zh_ji	Locale/runtime result	Decision
			_micro	eba_micro		
Whisper large-v3	Fixed 15-row contract	15	33.33	42.04	2 locale-violation rows, 23 simplified chars	Do not promote to 258-row or selected-300
Whisper large-v3-turbo	Fixed 15-row contract	15	40.36	50.87	4 locale-violation rows, 48 simplified chars	Retain as bounded feasibility evidence only; no split-level or selected-300 claim
SenseVoiceSmall	Fixed 15-row contract	15	63.12	78.98	14 locale-violation rows, 209 simplified chars	Reject from full split until locale policy changes
Qwen3-ASR-0.6B	Fixed 15-row contract	15	64.16	81.07	15 locale-violation rows, 260 simplified chars	Reject from full split until locale policy changes
Qwen3-ASR-1.7B	Bounded load gate	0	n/a	n/a	Timeout before inference at fetch/load after about 60.07s	Retry only after isolated cache/download plan
Gemma 4 E2B/E4B	Local multi-modal class probe	0	n/a	n/a	Local Transformers 4.57.6 has no required Gemma 4 multimodal/audio class	Build isolated prompted multimodal runtime before testing

Evidence sources: 70_experiments/runs/asr_candidate_current_recheck_2026_05_26/candidate_current_recheck_2026_05_26 and 70_experiments/runs/asr_candidate_current_recheck_2026_05_26/summary.json.

Whisper large-v3, Whisper large-v3-turbo, SenseVoiceSmall, and Qwen3-ASR-0.6B have field-contract evidence but fail strict Taiwan Traditional Chinese locale promotion criteria. Qwen3-ASR-1.7B remains stopped before inference at fetch/load, and Gemma 4 E2B/E4B remain blocked by local multimodal runtime support. The next validation layer is locale/runtime promotion, not a broader full-split ASR run.

Table 3. Human-Reviewed Predictor Table

Table 3 reports aggregate predictor comparison against human-reviewed `human_decision_change_yes` labels over 90 reviewed model assessments. WER/CER are transcript-surface baselines; SRES is the semantic-risk baseline; CEIS is the proposed decision-stability metric.

Target: `human_decision_change_yes`, 90 reviewed model assessments, 16 positive model assessments, clustered within 30 reviewed audio rows. Diagnostic thresholds are selected on the scoped audit set for aggregate comparison and are not frozen deployment thresholds.

Predictor	Unit	Row-clustered AUC		Diagnostic threshold	Best F1	Row-clustered F1 95% CI		Precision	Recall	False negative	Paper use
		AUC	95% CI								
WER	90 model assessments / 30 rows	0.6964	0.5699-0.8125	42.59	0.4615	0.2857-0.6667	0.3913	0.5625		7	Surface baseline
CER	90 model assessments / 30 rows	0.7276	0.6084-0.8372	16.42	0.4516	0.3158-0.6667	0.3043	0.8750		2	Surface baseline
SRES total	90 model assessments / 30 rows	0.8995	0.8119-0.9657	270.0	0.6512	0.4667-0.8421	0.5185	0.8750		2	Semantic-risk baseline
CEIS max	90 model assessments / 30 rows	0.9117	0.8516-0.9615	5.0	0.6275	0.4681-0.8148	0.4571	1.0000		0	Proposed decision-stability metric

Evidence source: 70_experiments/runs/janus_300_high_stakes_human_audit_selection_2026_05_25/hu
Row-clustered uncertainty source: 70_experiments/runs/janus_300_high_stakes_human_audit_selection
Leave-one-row-out sensitivity source: 70_experiments/runs/janus_300_high_stakes_human_audit_select

CEIS has the strongest point-estimate human-reviewed decision-change AUC and reaches recall 1.0 at the diagnostic threshold, while SRES achieves the highest best-threshold F1 and fewer false positives. Row-clustered AUC intervals overlap for CEIS and SRES, so the evidence supports CEIS as a conservative decision-stability signal rather than a universally dominant classifier.

Table 4. Human-Reviewed Recovery Policy Table

Table 4 reports five recovery policies evaluated as aggregate replay under human-reviewed selected-300 labels. The recovery evidence is aggregate-only and uses the same 30 reviewed rows and 90 reviewed model assessments as the predictor table.

Evidence mode: human-reviewed selected-300, 30 reviewed rows and 90 reviewed model assessments.

Policy	Model assessments	Unsafe down-routing	High-risk missed	Critical miss	Recovery budget	Row-clustered budget 95% CI	Severe misses eliminated	Row-clustered severe-miss 95% CI	Triggers per severe miss eliminated	Scoped interpretation
No recovery	90	29	6	1	0.0000	0.0000-0.0000	0	2-13 remaining severe misses	n/a	Baseline decision risk
Calibrated-confidence unavailable	90	29	6	1	0.0000	0.0000-0.0000	0	2-13 remaining severe misses	n/a	No calibrated confidence trigger available
SRES-triggered recovery	90	24	0	0	0.3889	0.3000-0.4889	7	0-0 remaining severe misses	5.0	Semantic-risk recovery baseline
CEIS-triggered conservative action	90	24	0	0	0.3889	0.3000-0.4889	7	0-0 remaining severe misses	5.0	Matches SRES severe-miss elimination at same budget in replay
CEIS ensemble arbitration	90	24	0	0	0.5222	0.4000-0.6556	7	0-0 remaining severe misses	6.7	Adds abstention/interval behavior at higher budget

Evidence sources: `70_experiments/runs/janus_300_high_stakes_recovery_human_reviewed_2026_05_26`, and `70_experiments/runs/janus_300_high_stakes_recovery_human_reviewed_2026_05_26/summary.json`.
Row-clustered uncertainty source: `70_experiments/runs/janus_300_high_stakes_recovery_human_reviewed_2026_05_26/row_clustered_uncertainty.json`.
Leave-one-row-out sensitivity source: `70_experiments/runs/janus_300_high_stakes_recovery_human_reviewed_2026_05_26/leave_one_row_out_sensitivity.json`.
Fixed-budget frontier source: `70_experiments/runs/janus_300_high_stakes_recovery_human_reviewed_2026_05_26/fixed_budget_frontier.json`.

Without recovery, the reviewed evidence contains 6 high-risk misses and 1 critical miss. SRES-triggered recovery and CEIS-triggered conservative action both eliminate high-risk missed and critical miss counts in aggregate policy replay at the same 0.3889 trigger budget, corresponding to 35 triggers for 7 severe missed outcomes eliminated. CEIS ensemble arbitration preserves this 0/0 result while introducing abstention behavior at a higher 0.5222 budget. The policy replay eliminates the most severe missed-risk outcomes under the scoped labels, while residual unsafe downrouting remains at 24 and requires separate governance.

A fixed-budget replay provides an additional operating view. At 10%, 20%, 30%, and 40% requested trigger budgets, CEIS-ranked conservative replay leaves 0 high-risk misses and 0 critical misses under the scoped labels; the 10% point uses 9 triggers. SRES-ranked conservative replay leaves 4, 2, 2, and 0 severe missed outcomes at the same requested budgets, reaching 0 severe misses only when all 35 eligible SRES triggers are used. This frontier supports CEIS as a conservative decision-stability signal while preserving the Table 4 result that the diagnostic SRES and CEIS policies tie at the selected 0.3889 budget. Appendix Table A1 reports the fixed-budget replay explicitly as an operating-point frontier rather than a deployment-threshold claim.

Figure Package

The figure package is generated from aggregate-only inputs by `80_semantic_risk_asr/paper/generate_paper_figures.py`. The Python entrypoint `80_semantic_risk_asr/paper/generate_paper_figures.py` remains only as a compatibility wrapper that calls the R script. The generated SVGs and PDF exports and their privacy boundary are listed in `80_semantic_risk_asr/paper/figures/`. The aggregate artifact manifest is generated by `80_semantic_risk_asr/paper/build_artifact_manifest.py` and written to `80_semantic_risk_asr/paper/artifact_manifest.tsv`.

Figure	File	Caption	Source	Privacy boundary
F1. CDS-ASR pipeline	figures/f1_cds_asr_pipeline.svg	CDS-ASR converts audio into ASR hypotheses and runtime signals, extracts risk atoms, generates plausible variants, scores SRES/CEIS, and applies constrained recovery or conservative action.	Method section	No transcript text or row content
F2. Evidence boundary	figures/f2_evidence_boundary.svg	The manuscript separates 258-row split/model-comparison evidence, selected-300 provenance evidence, and selected-300 human-reviewed predictor/recovery evidence.	publishable evidence summary	Aggregate status only
F3. Predictor AUC	figures/f3_predictor_auc.svg	CEIS has the highest human-reviewed decision-change AUC in the scoped selected-300 audit, while SRES remains strongest on best-threshold F1; Table 3 reports row-clustered uncertainty.	human_audit_predictor_comparison.tsv	Aggregate predictor metrics
F4. Recovery outcomes	figures/f4_recovery_outcomes.svg	SRES-triggered recovery and CEIS-triggered conservative action both eliminate high-risk missed and critical miss counts in aggregate replay at the same 0.3889 budget.	policy_comparison.tsv	Aggregate policy counts
F5. Model lane state	figures/f5_model_lane_state.svg	Main comparable split evidence, locale-gated candidates, and runtime-blocked probes remain separate until promotion gates are satisfied.	main/candidate aggregate summaries	Aggregate lane state
F6. N-ladder	figures/f6_n_ladder.svg	The evidence chain separates 258 rows, selected-300 provenance, 30 reviewed rows, and 90 clustered model assessments.	method evidence units	Aggregate counts only
F7. Budget-risk frontier	figures/f7_budget_risk_frontier.svg	Fixed-budget policy replay shows the trigger-budget trade-off needed to eliminate severe missed outcomes under scoped labels.	fixed_budget_recovery_frontier.tsv	Aggregate policy counts
F8. Low-WER danger signals	figures/f8_low_wer_danger.svg	Low-WER rows can still carry label-flip, unsafe-downrouting, and conservative-trigger signals, directly supporting the title claim that low WER can remain operationally dangerous.	low_wer_danger_summary.tsv	Aggregate proxy counts
F9. Risk-atom instability heatmap	figures/f9_risk_atom_instability_heatmap.svg	Instability concentrates differently by ASR model and decision-critical atom, showing the mechanism behind CEIS/SRES risk signals.	risk_atom_instability.tsv	Aggregate proxy counts
F10. Human-reviewed risk-atom outcomes	figures/f10_human_reviewed_atom_outcomes.svg	Human review connects risk-atom coverage to criticality and decision-change evidence across reviewed rows.	human_audit_risk_atom_review.tsv	Aggregate human-reviewed counts
F11. Risk-atom entropy heatmap	figures/f11_risk_atom_entropy_heatmap.svg	Atom-class entropy summarizes whether instability is concentrated in a few atom classes or distributed across several atom classes.	risk_atom_instability.tsv	Aggregate entropy from atom counts

Submission-package PDF exports are generated with the same basename under

80_semantic_risk_asr/paper/figures/.

Word-level confusion entropy is a valuable next analysis, especially for frequently misrecognized decision words around negation, amount, transfer action, actor identity, and scam-pattern cues. The current submission package keeps that analysis out of the public figure set unless it is first converted into a redacted aggregate lexicon with no transcript spans, audio IDs, selected row IDs, hypotheses, or reviewer notes. The R-generated F11 is therefore the publishable entropy layer for the current evidence boundary, while a future local-only confusion lexicon can support a supplemental word-level entropy heatmap after privacy review.

Discussion

The evidence supports a consequence-centered ASR evaluation claim: correct WER/CER reporting is necessary, but transcript similarity alone does not capture high-stakes decision instability. The reviewed selected-300 predictor table shows that SRES and CEIS align more strongly with human-reviewed decision-change labels than WER/CER in this scoped evidence layer.

The model comparison should be presented as an evidence layer, not as a model leaderboard. The partial encoder is the strongest current ASR hypothesis generator on the comparable split, and the selected-300 human-reviewed evidence shows how downstream decision evidence refines transcript-centered model comparison.

The recovery result supports a scoped replay claim. Under human-reviewed selected-300 labels, risk-triggered conservative policies, including CEIS-triggered conservative action, eliminate high-risk missed and critical miss counts in aggregate replay. SRES-triggered recovery and CEIS-triggered conservative action reach this result at the same 0.3889 trigger budget, while ensemble arbitration adds interval/abstention behavior as a higher-budget governance option. CEIS is therefore best presented as the more conservative decision-instability signal at the predictor layer, while SRES remains a strong semantic-risk recovery baseline.

Improving ASR remains necessary, but stronger transcript models do not replace decision-stability analysis. The paper studies residual instability after transcript scoring: whether plausible ASR alternatives around decision atoms would change escalation, routing, or conservative action. This makes CDS-ASR a safety layer for the remaining decision interval, not a substitute for ASR model improvement.

The aggregate-only artifact boundary is also a contribution to reproducible governance. The paper does not claim full public row-level reproducibility. Instead, it provides aggregate reproducibility, operation records, validation gates, and consistency checks under a privacy-preserving audit boundary.

The candidate lane is already bounded by evidence. Whisper large-v3, Whisper large-v3-turbo, SenseVoiceSmall, and Qwen3-ASR-0.6B have small-gate evidence but fail strict Taiwan Traditional Chinese locale promotion criteria. Qwen3-ASR-1.7B is stopped by fetch/load timeout, and Gemma 4 E2B/E4B require an isolated multimodal runtime. The next validation step is locale/runtime validation, not a broader full-split ASR run.

Claim Registry

Claim	Scope	Evidence artifact	Statistic	Limitation / defense
CEIS has the strongest AUC among reported predictors	90 clustered model assessments from 30 reviewed rows	<code>human_audit_predict_or_comparison.tsv</code>	AUC 0.9117	Row-clustered CI needed; selected high-stakes audit surface
CEIS reaches a zero-FN operating point	scoped selected-300 diagnostic threshold	<code>human_audit_predict_or_comparison.tsv</code>	FN 0, recall 1.0000	Retrospective diagnostic threshold, not frozen deployment threshold
SRES and CEIS conservative policies eliminate severe misses in replay	aggregate policy replay over reviewed assessments	<code>policy_comparison.tsv</code>	high-risk missed 0, critical miss 0	Replay evidence, not live causal deployment; policies tie at 0.3889 budget
Partial encoder is the strongest current ASR hypothesis generator	258-row split/model-comparison layer	<code>asr_cds_proxy_comparison.tsv</code>	<code>cer_zh_micro</code> 15.04, unsafe downrouting 7, high-risk missed 4	Split-level model-comparison context only
Selected-300 is a high-stakes audit surface	enriched selection provenance	<code>human_audit_selection_summary.json</code> , <code>selection_strata.tsv</code>	300 candidates, 30 reviewed rows, 90 assessments	Not prevalence-preserving; selection enrichment disclosed
Aggregate-only release supports reviewer-visible auditability	release and operation-record layer	validation summaries, evidence matrices, operation records, figure scripts	gate state: roadmap complete, publishable ready, consistency 26/26	No public row-level transcript reproducibility

Ethics, Privacy, And Intended Use

This study treats transcript-bearing speech data as sensitive operational content. NIH human-subjects guidance treats the use, study, analysis, or generation of identifiable private information as a human-subjects trigger under the cited federal definition (National Institutes of Health 2024). This paper therefore keeps raw audio, transcript-bearing hypotheses, reviewer notes, row identifiers, local response sheets, and runtime logs outside the release boundary. Before any external data release or deployment claim, the institutional review, exemption, data-use, retention, encryption, and access control status must be documented.

CDS-ASR is intended for ASR safety audit, risk-aware routing, manual-review

prioritization, conservative escalation, and machine abstention. It is not intended for automatic guilt or fraud determination, automatic account freezing, service denial, punitive action, or direct law-enforcement reporting without human review. “Conservative machine action” means preserving uncertainty, raising review priority, abstaining, or requiring human confirmation.

This intended-use boundary aligns the paper with risk-management governance rather than autonomous adverse decision-making. NIST describes the AI RMF as a voluntary framework for managing AI risks to individuals, organizations, and society and for incorporating trustworthiness considerations into AI design, development, use, and evaluation (National Institute of Standards and Technology 2026). For cross-border or future deployment contexts, the EU AI Act illustrates the broader regulatory direction toward risk-based obligations and high-risk AI governance (European Commission 2026). This manuscript does not provide legal compliance analysis; it provides a scoped consequence-centered audit method and a privacy-preserving release boundary.

Limitations And Threats To Validity

The human-reviewed evidence is scoped to 30 rows and 90 model assessments. It supports a focused high-stakes audit claim, not a population-level deployment claim.

The 90 model assessments are clustered within 30 audio rows and should not be treated as 90 independent calls. The current submission-prep package reports row-clustered bootstrap intervals and leave-one-row-out sensitivity tables for predictor and recovery metrics; these remain uncertainty descriptions for a scoped audit rather than population-level deployment intervals.

The number of positive decision-change cases is limited. Table 3 has 16 positive model assessments, so AUC and threshold behavior should be interpreted with uncertainty.

The selected-300 audit surface is deliberately enriched for high-stakes signals. It supports decision-stability behavior within the selected audit boundary, not population prevalence of ASR-induced harm across all calls.

The reported best thresholds are diagnostic operating points on the scoped reviewed audit. They can overstate deployment performance unless threshold selection is later frozen on a separate development set.

Human labels come from a completed expert audit rather than multi-annotator adjudication. No second-reviewer blinded spot-check is included in the current submission-prep package, so inter-annotator agreement is not claimed. This is a single-expert audit boundary, not a multi-annotator reliability claim.

Recovery evidence is aggregate-only. This protects transcript-bearing call and review content, but limits external row-level reproducibility.

CEIS depends on generated plausible variants and risk atom weights. Missed variants or misweighted atoms can affect instability scoring. The current aggregate package records risk-atom coverage and a variant-coverage status table, but a full source-specific aggregate variant generation audit remains a planned validation extension.

The Taiwan Traditional Chinese locale policy is strict by design. Candidate model rejection may reflect deployment-locale mismatch, not universal ASR inferiority.

Recovery policies eliminate high-risk missed and critical miss counts in aggregate replay, but Table 4 still leaves unsafe downrouting count at 24 after SRES/CEIS conservative recovery. The evidence supports scoped severe-miss reduction, not elimination of all safety risk.

The study does not include a deployment trial and does not estimate population prevalence of ASR-induced harm across all calls. It evaluates whether decision-critical ASR risk can be detected and mitigated within a deliberately enriched high-stakes audit surface.

Appendix / Artifact Availability

Artifact Availability Statement

The repository release includes aggregate run records, validation summaries, metric tables, evidence matrices, figure-generation code, and paper-facing documentation. Raw audio, raw transcripts, selected row identifiers, audio identifiers, model hypotheses, transcript-bearing runtime logs, reviewer response sheets, reviewer notes, and model weights are not released because they may contain sensitive call or review content.

Because transcript-bearing materials may contain sensitive call content, reproducibility is provided through aggregate validation artifacts, operation records, manifest checks, and consistency audits rather than through raw row release. This provides reviewer-visible auditability for model comparison, metric policy, selected-300 human-review status, predictor analysis, recovery-policy evaluation, candidate-lane boundaries, and consistency checks while preserving the local-only boundary for sensitive materials.

The reviewer package should include an aggregate manifest with artifact path, role, privacy class, generator, source inputs, SHA256, source git commit, timestamp, Python version, metric-library versions, tokenizer policy, locale normalization policy, decoding parameters where available, random seeds where used, and hardware summary. The manifest records auditability metadata without adding transcript-bearing row content.

Reviewer-Reproducible Aggregate Artifacts

- Artifact manifest: 80_semantic_risk_asr/paper/artifact_manifest.tsv
- Artifact manifest generator: 80_semantic_risk_asr/paper/build_artifact_manifest.py
- Experiment registry: 70_experiments/registry.tsv
- Main 258-row comparison: 70_experiments/runs/janus_258_test_split_asr_cds_proxy/asr_cds_proxy
- Metric-definition audit: 70_experiments/runs/wer_metric_audit_2026_05_25/journal_compliance
- Selected-300 human audit refresh: 70_experiments/runs/janus_300_high_stakes_human_audit_selected
- Human-reviewed predictor evidence: 70_experiments/runs/janus_300_high_stakes_human_audit_selected
- Human-reviewed predictor row-clustered CI: 70_experiments/runs/janus_300_high_stakes_human_audit_selected
- Human-reviewed predictor leave-one-row-out sensitivity: 70_experiments/runs/janus_300_high_stakes_human_audit_selected
- Human-reviewed recovery evidence: 70_experiments/runs/janus_300_high_stakes_recovery_human_audit_selected
- Human-reviewed recovery row-clustered CI: 70_experiments/runs/janus_300_high_stakes_recovery_human_audit_selected
- Human-reviewed recovery leave-one-row-out sensitivity: 70_experiments/runs/janus_300_high_stakes_recovery_human_audit_selected
- Human-reviewed recovery fixed-budget frontier: 70_experiments/runs/janus_300_high_stakes_recovery_human_audit_selected
- Claim registry: 70_experiments/runs/postdoc_evidence_chain_2026_05_25/claim_registry.tsv
- CEIS method summary: 70_experiments/runs/postdoc_evidence_chain_2026_05_25/ceis_method_summary
- Selection provenance summary: 70_experiments/runs/janus_300_high_stakes_human_audit_selected
- Counterfactual variant coverage status: 70_experiments/runs/janus_300_high_stakes_human_audit_selected
- Post-review evidence checklist: 70_experiments/runs/janus_300_high_stakes_human_audit_selected
- Publishable evidence audit: 70_experiments/runs/postdoc_evidence_chain_2026_05_25/publishable_evidence_audit
- Consequence evidence matrix: 70_experiments/runs/postdoc_evidence_chain_2026_05_25/consequence_evidence_matrix
- Roadmap completion audit: 70_experiments/runs/postdoc_evidence_chain_2026_05_25/postdoc_completion_audit
- Evidence-chain consistency audit: 70_experiments/runs/postdoc_evidence_chain_2026_05_25/evidence_chain_consistency_audit
- Candidate bounded recheck: 70_experiments/runs/asr_candidate_current_recheck_2026_05_26/su

Appendix Table A1. Fixed-Budget Recovery Frontier

The fixed-budget frontier is a retrospective aggregate replay over the same 90 reviewed model assessments. It reports what remains after applying ranked conservative triggers at requested trigger budgets. Requested 40% budget maps to 35 eligible triggers, so the observed budget is 0.3889.

Score metric	Requested budget	Trigger	Observed budget	Unsafe down- routing	High- risk missed	Critical miss	Severe misses	Severe misses eliminated	Triggers per severe miss eliminated
SRES total	0.1000	9	0.1000	28	4	0	4	3	3.0
SRES total	0.2000	18	0.2000	26	2	0	2	5	3.6
SRES total	0.3000	27	0.3000	26	2	0	2	5	5.4
SRES total	0.4000	35	0.3889	24	0	0	0	7	5.0
CEIS	0.1000	9	0.1000	24	0	0	0	7	1.2857
CEIS	0.2000	18	0.2000	24	0	0	0	7	2.5714
CEIS	0.3000	27	0.3000	24	0	0	0	7	3.8571
CEIS	0.4000	35	0.3889	24	0	0	0	7	5.0

Evidence source: 70_experiments/runs/janus_300_high_stakes_recovery_human_reviewed_2026_05_26/su

Operation Records

The selected-300 reviewer route is documented through aggregate-only operation records, including response closeout, response apply logs, reviewer work order, post-review sequence, operation-record audit, and consistency audit. These records provide command-level reproducibility without exposing transcript or row content.

Key records:

- 70_experiments/runs/janus_300_high_stakes_human_audit_selection_2026_05_25/human_audit_
- 70_experiments/runs/janus_300_high_stakes_human_audit_selection_2026_05_25/human_audit_
- 70_experiments/runs/janus_300_high_stakes_human_audit_selection_2026_05_25/human_audit_
- 70_experiments/runs/janus_300_high_stakes_human_audit_selection_2026_05_25/human_audit_
- 70_experiments/runs/janus_300_high_stakes_human_audit_selection_2026_05_25/human_audit_

Privacy And Local-Only Boundary

The following stay local or ignored:

- raw audio;
- raw transcripts;
- selected sample IDs and audio IDs;
- transcript-bearing audit sheets and response sheets;
- reviewer notes;
- raw model hypotheses and predictions;
- transcript-bearing runtime logs;
- checkpoints, adapters, and model weights;
- local downloaded reviewer packets and zip files.

The tracked repo preserves aggregate counts, run records, validation summaries, metric tables, and paper-facing evidence matrices. This boundary lets reviewers audit the evidence chain while protecting transcript-bearing material.

Validation Gate Commands

Run these after manuscript packaging changes:

```
.venv/bin/python 80_semantic_risk_asr/scoring/audit_postdoc_roadmap_completion.py --output-  
.venv/bin/python 80_semantic_risk_asr/scoring/audit_publishable_evidence_chain.py --output-  
.venv/bin/python 80_semantic_risk_asr/scoring/audit_evidence_chain_consistency.py --output-
```

Expected current gate state:

- roadmap completion: `roadmap_complete=true`, `blocking_gate=None`;
- publishable evidence: `publishable_ready=true`;
- consistency audit: 26/26 checks pass, `failed_checks=[]`.

Scope Control For Additional Experiments

Do not run new full-split ASR experiments now. Candidate models can move to 258-row or selected-300 only after strict Taiwan Traditional Chinese locale policy is satisfied or an isolated Gemma 4 multimodal runtime exists. Until then, the active work is manuscript drafting, results-table packaging, artifact availability, citation completion, and submission readiness.

- Amazon Web Services. 2026a. “Analyze Conversations Using Conversational Analytics in Connect Customer Contact Lens.” <https://docs.aws.amazon.com/connect/latest/adminguide/analyze-conversations.html>.
- . 2026b. “Contact Center Conversational Analytics - Amazon Connect.” <https://aws.amazon.com/products/connect/customer/conversational-analytics/>.
- Angelopoulos, Anastasios N., and Stephen Bates. 2021. “A Gentle Introduction to Conformal Prediction and Distribution-Free Uncertainty Quantification.” <https://arxiv.org/abs/2107.07511>.
- Chow, C. K. 1970. “On Optimum Recognition Error and Reject Tradeoff.” *IEEE Transactions on Information Theory* 16 (1): 41–46. <https://doi.org/10.1109/TIT.1970.1054406>.
- European Commission. 2026. “AI Act.” <https://digital-strategy.ec.europa.eu/en/policies/regulatory-framework-ai>.
- Federal Bureau of Investigation. 2026a. “2025 Internet Crime Report.” https://www.fbi.gov/file-repository/2025_ic3report.pdf.
- . 2026b. “Cryptocurrency and AI Scams Bilk Americans of Billions.” <https://www.fbi.gov/news/press-releases/cryptocurrency-and-ai-scams-bilk-americans-of-billions>.
- Geifman, Yonatan, and Ran El-Yaniv. 2017. “Selective Classification for Deep Neural Networks.” In *Advances in Neural Information Processing Systems 30*. <https://papers.neurips.cc/paper/7073-selective-classification-for-deep-neural-networks>.
- . 2019. “SelectiveNet: A Deep Neural Network with an Integrated Reject Option.” In *Proceedings of the 36th International Conference on Machine Learning*, 97:2151–59. Proceedings of Machine Learning Research. PMLR. <https://proceedings.mlr.press/v97/geifman19a.html>.
- Kim, Suyoun, Abhinav Arora, Duc Le, Ching-Feng Yeh, Christian Fuegen, Ozlem Kalinli, and Michael L. Seltzer. 2021. “Semantic Distance: A New Metric for ASR Performance Analysis Towards Spoken Language Understanding.” In *Interspeech 2021*, 1977–81. <https://doi.org/10.21437/Interspeech.2021-1929>.
- Miner, Adam S., Albert Haque, Jason A. Fries, et al. 2020. “Assessing the Accuracy of Automatic Speech Recognition for Psychotherapy.” *Npj Digital Medicine* 3: 82. <https://doi.org/10.1038/s41746-020-0285-8>.
- Naderi, Maryam, Enno Hermann, Alexandre Nanchen, Sevada Hovsepyan, and Mathew Magimai.-Doss. 2024. “Towards Interfacing Large Language Models with ASR Systems Using Confidence Measures and Prompting.” In *Interspeech 2024*, 2980–84. <https://doi.org/10.21437/Interspeech.2024-989>.

- National Institute of Standards and Technology. 2026. “AI Risk Management Framework.” <https://www.nist.gov/itl/ai-risk-management-framework>.
- National Institutes of Health. 2024. “Definition of Human Subjects Research.” <https://grants.nih.gov/policy-and-compliance/policy-topics/human-subjects/research>.
- National Police Agency, Ministry of the Interior. 2024. “165 Anti-Fraud Hotline.” <https://www.npa.gov.tw/en/app/artwebsite/view?id=8035&module=artwebsite&serno=ed2427e1-de0a-4f6f-8f68-8f83b604e89b>.
- Rugayan, Janine, Giampiero Salvi, and Torbjørn Svendsen. 2023. “Perceptual and Task-Oriented Assessment of a Semantic Metric for ASR Evaluation.” In *Interspeech 2023*, 2158–62. <https://doi.org/10.21437/Interspeech.2023-1778>.